

# Курс «Глубокое обучение в компьютерном зрении»

## Урок 3. Продвинутые архитектуры свёрточных нейросетей

### Практическое задание 3

Обучить СНС с помощью Transfer Learning на датасете Food-101  
Использовать тонкую настройку существующей предобученной модели и методы аугментации данных.  
Библиотеки: [Python, Tensorflow]

Выполнил **Соковнин ИЛ**

### Food-101

<https://www.tensorflow.org/datasets/catalog/food101>

**Описание :**

Этот набор данных состоит из 101 категории продуктов питания и 101 000 изображений. Для каждого класса предоставляется 250 проверенных вручную тестовых изображений, а также 750 обучающих изображений. Учебные изображения намеренно не очищались и поэтому все еще содержат некоторое количество шума. В основном это проявляется в виде интенсивных цветов и иногда неправильных этикеток. Все изображения были масштабированы таким образом, чтобы максимальная длина стороны составляла 512 пикселей.

- Домашняя страница : [https://data.vision.ee.ethz.ch/cvl/datasets\\_extra/food-101/](https://data.vision.ee.ethz.ch/cvl/datasets_extra/food-101/)

```
%tensorflow_version 2.x

%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np

import tensorflow as tf
import tensorflow_datasets as tfds

from skimage.transform import resize
from tensorflow.keras import layers

gpu_devices = tf.config.experimental.list_physical_devices('GPU')
for device in gpu_devices:
    tf.config.experimental.set_memory_growth(device, True)
```

### Загрузка датасета Food-101

Сохранение... ✕

```
tfds.disable_progress_bar()
(train_ds, test_ds), ds_info = tfds.load(
    'food101',
    as_supervised=True,
    with_info=True,
    split=['train[:90%]', 'train[90%:]'], # ?
)

Downloading and preparing dataset food101/2.0.0 (download: 4.65 GiB, generated: Unknown size, total: 4.65 GiB) to /root/tensorf
Shuffling and writing examples to /root/tensorflow_datasets/food101/2.0.0.incomplete4D5QJZ/food101-train.tfrecord
Shuffling and writing examples to /root/tensorflow_datasets/food101/2.0.0.incomplete4D5QJZ/food101-validation.tfrecord
Dataset food101 downloaded and prepared to /root/tensorflow_datasets/food101/2.0.0. Subsequent calls will reuse this data.
CPU times: user 2min 4s, sys: 44.9 s, total: 2min 49s
Wall time: 8min 7s
```

```
ds_info.description
```

```
'This dataset consists of 101 food categories, with 101'000 images. For each class, 250 manually reviewed test images are provided as well as 750 training images. On purpose, the training images were not cleaned, and thus still contain some amount of noise. This comes mostly in the form of intense colors and sometimes wrong labels. All images were rescaled to have a maximum side length of 512 pixels.'
```

```
# Get class names
```

```
classes = ds_info.features["label"].names
```

```
classes[:10]
```

```
['apple_pie',  
'baby_back_ribs',  
'baklava',  
'beef_carpaccio',  
'beef_tartare',  
'beet_salad',  
'beignets',  
'bibimbap',  
'bread_pudding',  
'breakfast_burrito']
```

```
number_classes = len(classes)
```

```
number_classes
```

```
101
```

```
print(type(train_ds))
```

```
<class 'tensorflow.python.data.ops.dataset_ops.PrefetchDataset'>
```

## ▼ Визуализация датасета Food0101

```
%%time
```

```
some_samples = []
```

```
some_labels = []
```

```
for x in iter(test_ds.take(32)):
```

```
    some_samples.append(x[0])
```

```
    some_labels.append(x[1])
```

```
fig = plt.figure(figsize=(16, 8))
```

```
for j in range(len(some_samples)):
```

```
    ax = fig.add_subplot(4, 8, j+1)
```

```
    ax.imshow(some_samples[j])
```

```
    ax.title.set_text(classes[some_labels[j]])
```

```
    plt.xticks([], plt.yticks([]))
```

```
plt.show()
```



## ▼ Создание пайплайна данных

```
# Так как мы будем использовать готовую архитектуру MobileNet с предобученными весами,
# они ожидают конкретные картинки, в конкретном диапазоне [-1, 1].

INP_SIZE = 224
NUM_EPOCHS = 5
BATCH_SIZE = 32

def prepare(img, label):
    # Делаем диапазоне [-1, 1]. Такие картинки нельзя визуализировать.
    # Создаём квадратные картинки с входным размером 224x224.
    img = tf.cast(img, tf.float32)/127. - 1. # [0, 2] - 1 = [-1, 1]
    return tf.image.resize(img, (INP_SIZE, INP_SIZE)), label

train_ds = train_ds.shuffle(buffer_size=1000)
train_ds = train_ds.map(prepare)
train_ds = train_ds.batch(BATCH_SIZE, drop_remainder=True)

test_ds = test_ds.shuffle(buffer_size=1000)
test_ds = test_ds.map(prepare)
test_ds = test_ds.batch(BATCH_SIZE, drop_remainder=True) # размер batcha=128, drop_remainder - ?

CPU times: user 1.9 s, sys: 105 ms, total: 2 s
Wall time: 2.25 s

train_ds, test_ds

(<BatchDataset element_spec=(TensorSpec(shape=(32, 224, 224, 3), dtype=tf.float32, name=None), TensorSpec(shape=(32,), dtype=tf.int64, name=None))>,
 <BatchDataset element_spec=(TensorSpec(shape=(32, 224, 224, 3), dtype=tf.float32, name=None), TensorSpec(shape=(32,), dtype=tf.int64, name=None))>)
```

## ▼ Transfer Learning с дообучением весов предобученной модели (Fine-tuning)

**Fine-Tuning** - тонкая настройка

## ▼ Подготовка модели CNN

# Регуляризация - наложение ограничений на обучаемые параметры NN (решение проблемы переобучения).

```
# Вводим для модели L2-регуляризацию (L2 также называется WEIGHT_DECAY)
# т.е. 2-ю норму весов модели стремим к 0
WEIGHT_DECAY = 0.001 # Коэффициент lambda при 2-м слагаемом нашего лосса
```

```
wd = tf.keras.regularizers.l2(WEIGHT_DECAY) # L2(WEIGHT_DECAY) регуляризация
```

```
# см. 3_1_Advanced_CNN.ipynb
```

```
EXP_NAME = 'transfer'
base_model = tf.keras.applications.MobileNetV2(
    input_shape=(INP_SIZE, INP_SIZE, 3),
    include_top=False, # Верни модель без последнего слоя (top-слоя).
    weights='imagenet') # base_model будут предобучены на датасете Imagenet.
```

Сохранение...

```
for layer in base_model.layers[:-5]:
    base_model.trainable = True # Fine-tuning весов слоёв предобученной модели
```

```
for layer in base_model.layers[-5:]:
    layer.kernel_regularizer=wd # применяем L2(WEIGHT_DECAY)-регуляризацию
```

```
data_augmentation = tf.keras.Sequential([
    layers.experimental.preprocessing.RandomFlip("horizontal_and_vertical"),
    layers.experimental.preprocessing.RandomRotation((-0.2, 0.2)),
])
```

```
model = tf.keras.Sequential([
    tf.keras.layers.InputLayer(input_shape=(INP_SIZE, INP_SIZE, 3)),
    data_augmentation,
    base_model, # На выходе модели тензор с пространственными разметками [7x7x1280 (1280 каналов)]
```

```
# Добавляем собственные слои
tf.keras.layers.GlobalAveragePooling2D(), # Сокращаем все пространственные измерения и получаем вектор (по каналам)
tf.keras.layers.Dense(256, activation='relu', kernel_regularizer=wd),
tf.keras.layers.Dropout(0.5),
tf.keras.layers.Dense(101, activation='softmax')
])
```

▼ Model Summary

```
model.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
sequential (Sequential)	(None, 224, 224, 3)	0
mobilenetv2_1.00_224 (Functional)	(None, 7, 7, 1280)	2257984
global_average_pooling2d (GlobalAveragePooling2D)	(None, 1280)	0
dense (Dense)	(None, 256)	327936
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 101)	25957

=====

Total params: 2,611,877  
Trainable params: 2,577,765  
Non-trainable params: 34,112

=====

▼ Подготовка к обучению

```
LEARNING_RATE = 0.0001
optimizer = tf.keras.optimizers.Adam(lr=LEARNING_RATE)

model.compile(optimizer=optimizer,
              loss='sparse_categorical_crossentropy', # многоклассовая классификация
              metrics=['accuracy'])

tensorboard_callback = tf.keras.callbacks.TensorBoard(
    log_dir='logs/'+EXP_NAME,
    write_graph=False,
    update_freq=100,
    profile_batch=0)

/usr/local/lib/python3.7/dist-packages/keras/optimizer_v2/adam.py:105: UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
super(Adam, self).__init__(name, **kwargs)
```

▼ Обучение модели

Сохранение...

✕

```
%%time

history = model.fit(
    train_ds,
    epochs=NUM_EPOCHS,
    validation_data=test_ds,
    callbacks=[tensorboard_callback])
```

Epoch 1/5

2130/2130 [=====] - 475s 214ms/step - loss: 3.4858 - accuracy: 0.2958 - val\_loss: 2.4899 - val\_accuracy: 0.2958

Epoch 2/5

2130/2130 [=====] - 456s 213ms/step - loss: 2.3752 - accuracy: 0.5078 - val\_loss: 2.1219 - val\_accuracy: 0.5078

Epoch 3/5

2130/2130 [=====] - 456s 214ms/step - loss: 2.0160 - accuracy: 0.5784 - val\_loss: 1.8927 - val\_accuracy: 0.5784

```
Epoch 4/5
2130/2130 [=====] - 460s 216ms/step - loss: 1.7984 - accuracy: 0.6201 - val_loss: 1.6964 - val_accuracy: 0.6201
Epoch 5/5
2130/2130 [=====] - 458s 214ms/step - loss: 1.6325 - accuracy: 0.6521 - val_loss: 1.6614 - val_accuracy: 0.6521
CPU times: user 57min 4s, sys: 1min 33s, total: 58min 38s
Wall time: 39min 41s
```

Оценка качества модели

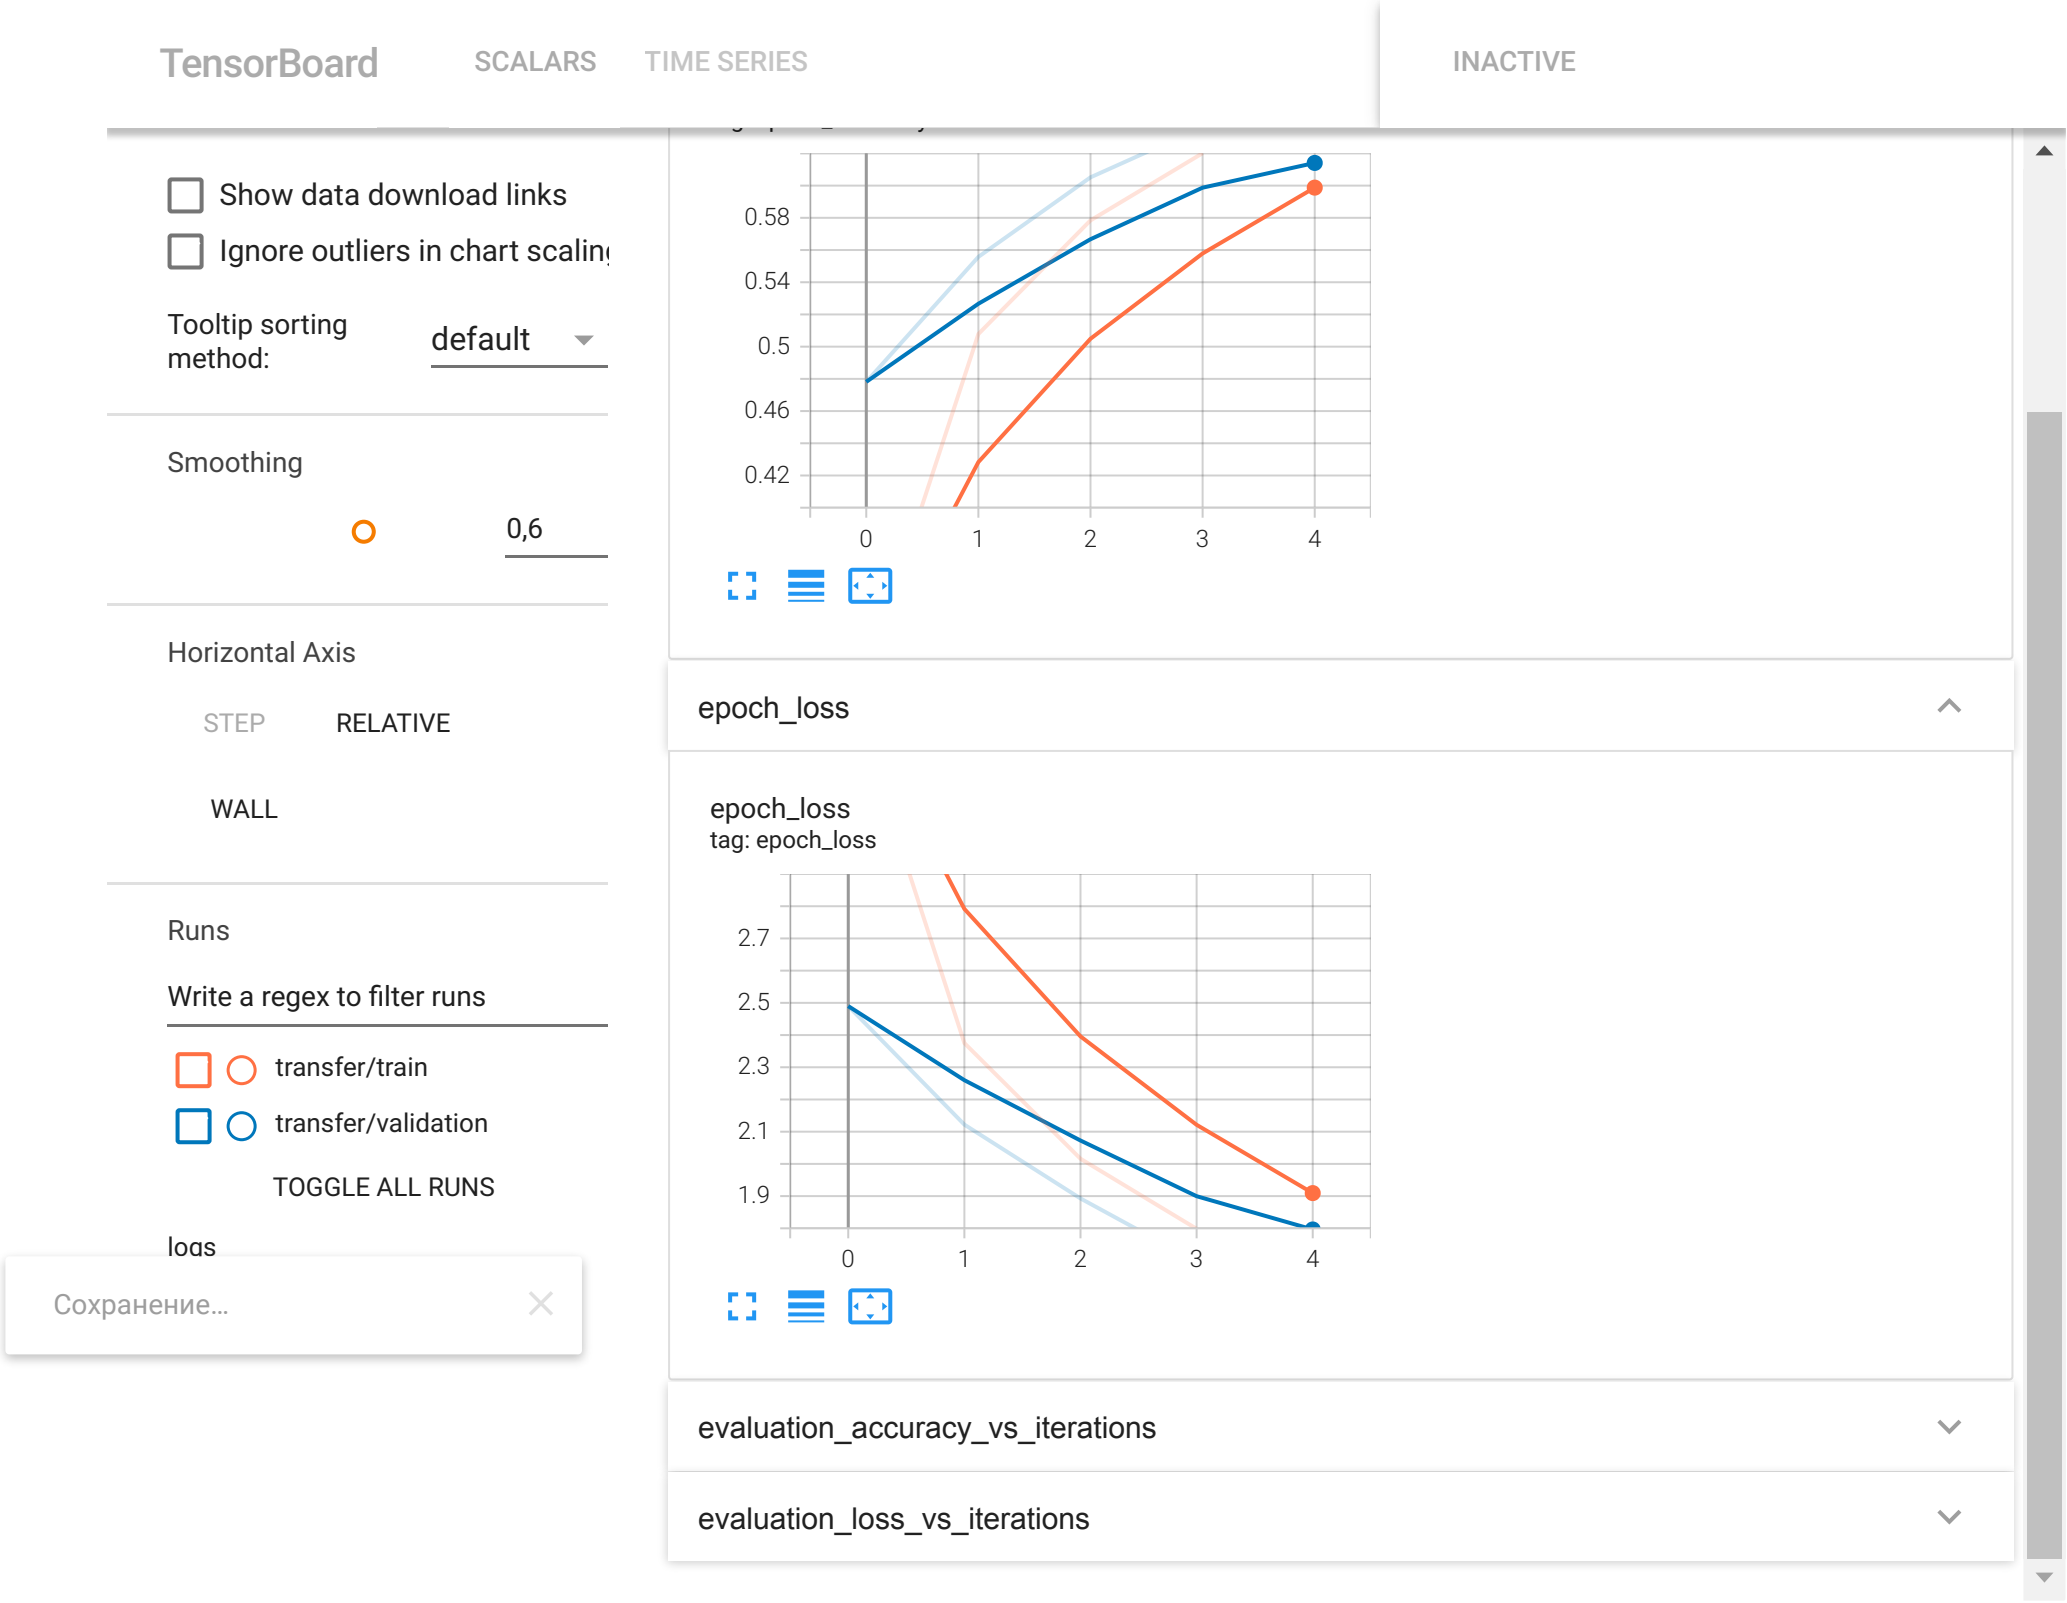
```
%%time

result = model.evaluate(test_ds)
result

236/236 [=====] - 25s 100ms/step - loss: 1.6602 - accuracy: 0.6344
CPU times: user 38.8 s, sys: 1.73 s, total: 40.6 s
Wall time: 41 s
```

TensorBoard

```
%load_ext tensorboard
%tensorboard --logdir logs
```



Запуск предсказания на изображениях

```
some_samples = []
some_labels = []
```

```
for x in iter(test_ds.take(32)):
    some_samples.append(x[0])
    some_labels.append(x[1])

predictions = model.predict(some_samples[0])

prediction_labels = tf.argmax(predictions, axis=1, output_type=tf.int32)

fig = plt.figure(figsize=(15, 30))
for j in range(len(some_samples[0])):
    ax = fig.add_subplot(8, 4, j+1)
    ax.imshow(some_samples[0][j])
    ax.title.set_text(f'pred: {classes[prediction_labels[j]]}\n' + \
                      f'true: {classes[some_labels[0][j]]}')
    plt.xticks([], plt.yticks([]))
plt.show()
```


Сохранение...





WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers)  
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers)  
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers)  
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers)  
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers)  
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers)

pred: pizza  
true: pizza



pred: beet\_salad  
true: beet\_salad




pred: onion\_rings  
true: onion\_rings



pred: pulled\_pork\_sandwich  
true: pulled\_pork\_sandwich



pred: macarons  
true: macarons



pred: paella  
true: scallops




pred: greek\_salad  
true: greek\_salad



pred: greek\_salad  
true: greek\_salad




pred: carrot\_cake  
true: baklava




pred: beef\_carpaccio  
true: beef\_carpaccio



pred: pulled\_pork\_sandwich  
true: pulled\_pork\_sandwich



pred: filet\_mignon  
true: filet\_mignon



pred: takoyaki  
true: falafel




pred: samosa  
true: gyoza



pred: pizza  
true: lasagna




pred: beef\_tartare  
true: beef\_tartare



pred: eggs\_benedict  
true: eggs\_benedict



pred: grilled\_cheese\_sandwich  
true: grilled\_cheese\_sandwich



pred: clam\_chowder  
true: lobster\_bisque



pred: french\_fries  
true: french\_fries




pred: red\_velvet\_cake

pred: creme\_brulee  
true: creme\_brulee




pred: chocolate\_mousse  
true: chocolate\_mousse



pred: pancakes  
true: french\_toast



pred: crab\_cakes  
true: crab\_cakes



pred: foie\_gras  
true: deviled\_eggs



pred: pulled\_pork\_sandwich  
true: pulled\_pork\_sandwich



pred: fried\_rice  
true: fried\_rice

